

Lecture 14

Collective Communication

Announcements

- Project Progress report, due next Weds 11/28

Today's lecture

- Collective Communication algorithms
- Sorting

Collective communication

- Collective operations are called by **all** processes within a communicator
- Basic collectives seen so far
 - ♦ Broadcast: distribute data from a designated root process to all the others
 - ♦ Reduce: combine data from all processes returning the result to the root process
 - ♦ Will revisit these
- Other Useful collectives
 - ♦ Scatter/gather
 - ♦ All to all
 - ♦ Allgather
- Diverse applications
 - ♦ Fast Fourier Transform
 - ♦ Sorting

Underlying assumptions

- Fast interconnect structure
 - ♦ All nodes are equidistant
 - ♦ Single-ported, bidirectional links
- Communication time is $\alpha + \beta n$ in the absence of contention
 - ♦ Determined by bandwidth β^{-1} for long messages
 - ♦ Dominated by latency α for short messages

Inside MPI-CH

- Tree like algorithm to broadcast the message to blocks of processes, and a linear algorithm to broadcast the message within each block
- Block size may be configured at installation time
- If there is hardware support (e.g. Blue Gene), then it is given responsibility to carry out the broadcast
- Polyalgorithms apply different algorithms to different cases, i.e. long vs. short messages, different machine configurations
- We'll use hypercube algorithms to simplify the special cases when $P=2^k$, k an integer

Details of the algorithms

- Broadcast
- AllReduce
- Scatter/gather
- Allgather
- All to all

Broadcast


- The root process transmits of m pieces of data to all the $p-1$ other processors
- Spanning tree algorithms are often used
- We'll look at a similar algorithm with logarithmic running time: the *hypercube algorithm*
- With the linear ring algorithm this processor performs $p-1$ sends of length m
 - ♦ Cost is $(p-1)(\alpha + \beta m)$

Sidebar: what is a hypercube?


- A hypercube is a d -dimensional graph with 2^d nodes
- A 0-cube is a single node, 1-cube is a line connecting two points, 2-cube is a square, etc
- Each node has d neighbors



1D





2D



3D


Properties of hypercubes

- A hypercube with p nodes has $\lg(p)$ dimensions
- *Inductive construction*: we may construct a d -cube from two $(d-1)$ dimensional cubes
- **Diameter**: What is the maximum distance between any 2 nodes?
- **Bisection bandwidth**: How many cut edges (mincut)




Bookkeeping

- Label nodes with a binary reflected grey code
<http://www.nist.gov/dads/HTML/graycode.html>
- Neighboring labels differ in exactly one bit position $001 = 101 \otimes e_2$, $e_2 = 100$



Hypercube broadcast algorithm with $p=4$

- Processor 0 is the root, sends its data to its hypercube “buddy” on processor 2 (10)
- Proc 0 & 2 send data to respective buddies



Details of the algorithms

- Broadcast
- AllReduce
- Scatter/gather
- Allgather
- All to all

Reduction


- We may use the hypercube algorithm to perform reductions as well as broadcasts
- Another variant of reduction provides all processes with a copy of the reduced result

Allreduce()

- Equivalent to a **Reduce** + **Bcast**
- A clever algorithm performs an **Allreduce** in one phase rather than having perform separate reduce and broadcast phases

Allreduce


- Can take advantage of duplex connections



Details of the algorithms

- Broadcast
- AllReduce
- Scatter/gather
- Allgather
- All to all

Scatter/Gather



Scatter

- Simple linear algorithm
 - ♦ Root processor sends a chunk of data to all others
 - ♦ Reasonable for long messages


$$(p-1)\alpha + \frac{p-1}{p}n\beta$$

- Similar approach taken for Reduce and Gather
- For short messages, we need to reduce the complexity of the latency (α) term

Minimum spanning tree algorithm

- Recursive hypercube-like algorithm with $\lceil \log P \rceil$ steps
 - ♦ Root sends half its data to process $(\text{root} + p/2) \bmod p$
 - ♦ Each receiver acts as a root for corresponding half of the processes
 - ♦ MST: organize communication along edges of a minimum-spanning tree covering the nodes
- Requires $O(n/2)$ temp buffer space on intermediate nodes
- Running time:

$$\lceil \lg P \rceil \alpha + \frac{p-1}{p} n \beta$$




Details of the algorithms


- Broadcast
- AllReduce
- Scatter/gather
- Allgather
- All to all

AllGather

- Equivalent to a gather followed by a broadcast
- All processors accumulate a chunk of data from all the others



AllGather



Allgather

- Use the all to all recursive doubling algorithm
- For P a power of two, running time is


$$\lceil \lg P \rceil \alpha + \frac{p-1}{p} n \beta$$

Details of the algorithms

- Broadcast
- AllReduce
- Scatter/gather
- Allgather
- All to all

All to all

- Also called *total exchange* or *personalized communication*: a transpose
- Each process sends a different chunk of data to each of the other processes
- Used in sorting and the Fast Fourier Transform



Exchange algorithm

- n elements / processor (n total elements)
- $p - 1$ step algorithm
 - ♦ Each processor exchanges n/p elements with each of the others
 - ♦ In step i , process k exchanges with processes $k \pm i$


for $i = 1$ to $p-1$

src = $(\text{rank} - i + p) \bmod p$

dest = $(\text{rank} + i) \bmod p$

sendrecv(from src to dest)

end for



- Good algorithm for long messages
- Running time:

$$(p-1)\alpha + (p-1)\frac{n}{p}\beta \approx n\beta$$

Recursive doubling for short messages

- In each of $\lceil \log p \rceil$ phases all nodes exchange $\frac{1}{2}$ their accumulated data with the others
- Only $P/2$ messages are sent at any one time

$D = 1$

while ($D < p$)

 Exchange & accumulate data with rank $\otimes D$

 Left shift D by 1

end while


- Optimal running time for short messages

$$\lceil \lg P \rceil \alpha + nP\beta \approx \lceil \lg P \rceil \alpha$$

Flow of information

A B C D

A B C D




A B C D

A B C D

P0	P1	P2	P3

Flow of information

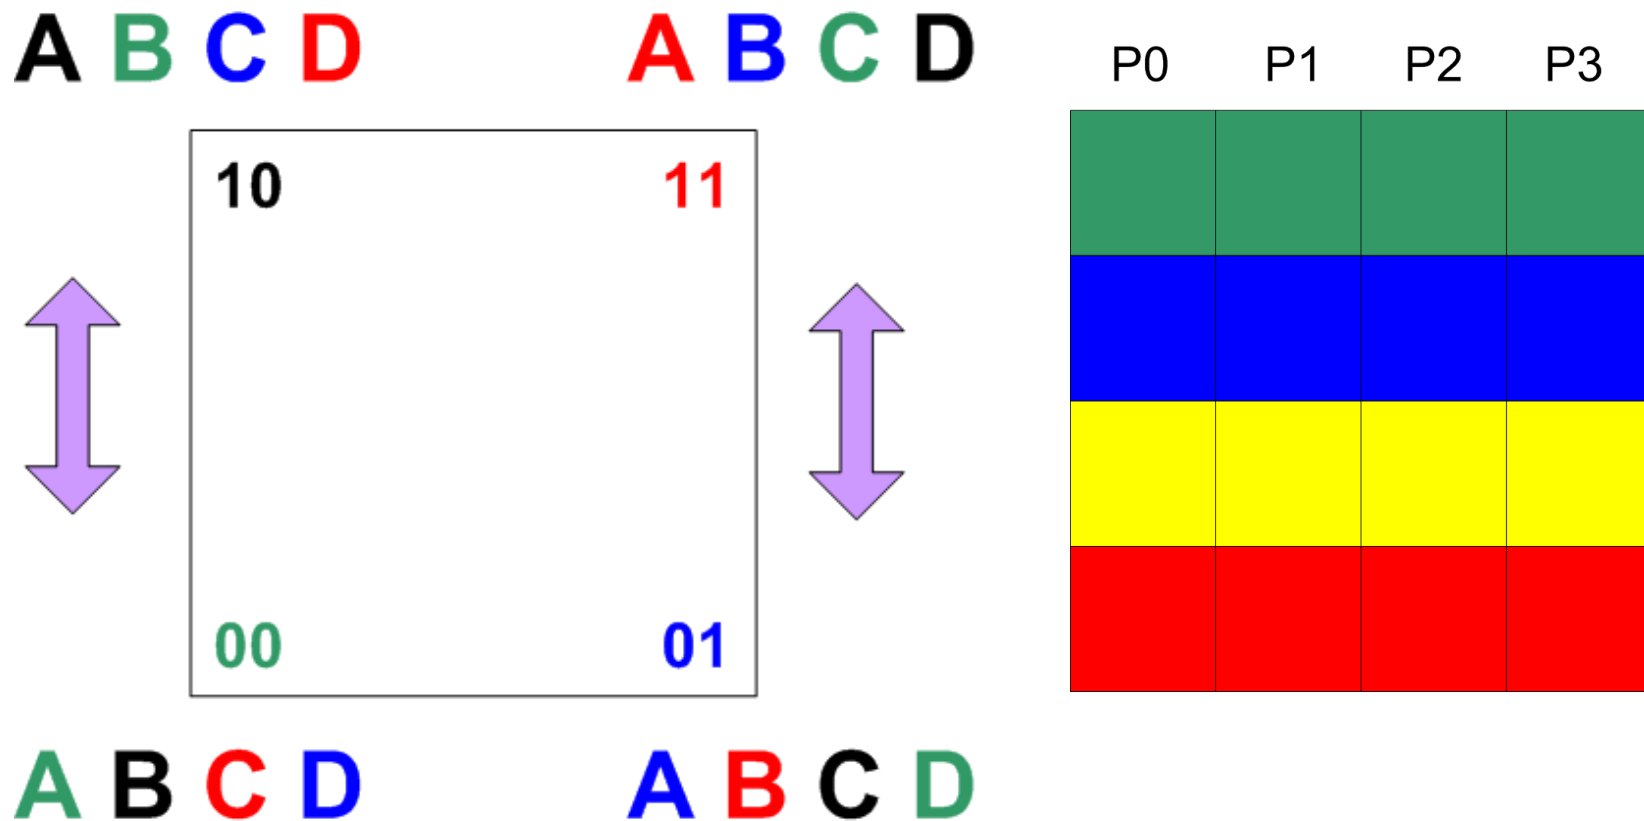
A B C D ↔ **A B C D**



A B C D ↔ **A B C D**

P0	P1	P2	P3
Green	Green	Yellow	Yellow
Blue	Blue	Red	Red
Green	Green	Yellow	Yellow
Blue	Blue	Red	Red

Flow of information



Summarizing all to all

- Short messages $\lceil \lg P \rceil \alpha$
- Long messages $\frac{P-1}{P} n \beta$


“Vector” All to All

- Generalize all-to-all, gather, etc.
- Processes supply varying length datum
- Vector all-to-all

`MPI_Alltoallv (`
 `void *sendbuf, int sendcounts[], int sDispl [],`
 `MPI_Datatype sendtype,`
 `void* recvbuf, int recvcnts[], int rDispl[],`
 `MPI_Datatype recvtype, MPI_Comm comm)`

- Used in sample sort (coming)

Alltoallv used in sample sort



Introduction to Parallel Computing, 2nd Ed., A. Grama, A. Gupta, G. Karypis, and V. Kumar, Addison-Wesley, 2003.

Details of the algorithms

- Broadcast
- AllReduce
- Scatter/gather
- Allgather
- All to all
- **Revisiting Broadcast**

Revisiting Broadcast

- P may not be a power of 2
- We use a binomial tree algorithm
- We'll use the hypercube algorithm to illustrate the special case of $P=2^k$
- Hypercube algorithm is efficient for short messages
- We use a different algorithm for long messages


Strategy for long messages

- Based van de Geijn's strategy
- Scatter the data
 - ♦ Divide the data to be broadcast into pieces, and fill the machine with the pieces
- Do an Allgather
 - ♦ Now that everyone has a part of the entire result, collect on all processors
- Faster than MST algorithm for long messages


$$2 \frac{p-1}{p} n\beta \ll \lceil \lg p \rceil n\beta$$

Algorithm for long messages

The scatter step



Algorithm for long messages



Today's lecture

- Collective Communication algorithms
- **Sorting**

Rank sorting

- Compute the rank of each input value
- Move each value in sorted position according to its rank
- On an ideal parallel computer, the **forall** loops parallelize perfectly



```
forall i=0:n-1, j=0:n-1
    if ( x[i] > x[j] ) then rank[i] += 1 end if
forall i=0:n-1
    y[rank[i]] = x[i]
```


In search of a fast and practical sort

- Rank sorting is impractical on real hardware
- Let's borrow the concept: compute the processor owner for each key
- Communicate data in sorted order in one step
- But how do we know which processor is the owner?
- Depends on the distribution of keys

Bucket sort

- Divide key space into equal subranges and associate a bucket with each subrange
- Unsorted input data distributed evenly over processors
- Each processor maintains p local buckets
 - ♦ Assigns each key to a local bucket: $\lfloor p \times \text{key} / (K_{\max} - 1) \rfloor$
 - ♦ Routes the buckets to the correct owner (each local bucket has $\sim n/p^2$ elements)
 - ♦ Sorts all incoming data into a single bucket




Wikipedia

Running time

- Assume that the keys are distributed uniformly over 0 to $K_{\max}-1$
- Local bucket assignment: $O(n/p)$
- Route each local bucket to the correct owner
All to all: $O(n)$
- Local sorting : $O(n/p)$
 - ◆ Radix sort
 - ◆ www.csse.monash.edu.au/~lloyd/tildeAlgDS/Sort/Radix

Scaling study

- IBM SP3 system: 16-way servers w/ Power 3 CPUs
- Weak scaling : 1M points per processor



Local sort: quicksort
 $O(n/p \log(n/p))$

All-to-all
 $O(n)$

Worst case behavior

- What is the worst case?
- Mapping of keys to processors based on knowledge of K_{\max}
- If keys are in range $[0, Q-1]$...
... processor k has keys in the range $[k*Q/P : (k+1)*Q/P]$
- For $Q=2^{30}$, $P=64$, each processor gets $2^{24} = 16$ M elements
- What if keys $\in [0, 2^{24}-1] \subset [0, 2^{30}-1]$?
- But if the keys are distributed non-uniformly, we need more information to ensure that the keys (and communication) are balanced over the processors
- Sample sort is an algorithm that collects such information and improves worst case behavior


Improving on bucket sort

- *Sample sort* remedies the problem
- “Parallel Sorting by Regular Sampling.” H. Shi and J. Schaeffer. *J. Parallel and Distributed Computing*, 14:361-372, 1992
- “Parallel Algorithms for Personalized Communication and Sorting With an Experimental Study.”
D. R. Helman, D.A. Bader, and J. JáJá,
Proc. SPAA: Annual ACM Symp. on Parallel Algorithms and Architectures (1996)
<http://www.umiacs.umd.edu/research/EXPAR/papers/spaa96.html>

The idea behind sample sort

- Use a heuristic to estimate the distribution of the global key range over the p processors processor so that...
- ...each processor gets about the same number of keys
- Sample the keys to determine a set of $p-1$ **splitters** that partition the key space into p disjoint intervals
[sample size parameter: s]
- Each interval is assigned a unique processor mapped to a bucket
- Once each processor knows the splitters, it can distribute its keys to the others accordingly
- Processors sort incoming keys

Alltoallv used in sample sort



Introduction to Parallel Computing, 2nd Ed., A.Grama, A.I Gupta, G. Karypis, and V. Kumar, Addison-Wesley, 2003.

Splitter selection: regular sampling

- After sorting local keys, each processor chooses p evenly spaced samples
- Each processor “deals” its sorted data into one of p bins
 - ♦ The k^{th} item is placed into position $\lceil k/p \rceil$ of bin $k \bmod p$
 - ♦ When done, each sends bin j to processor j
- This is like a transpose with block sizes $= n/p^2$
- Each processor receives **p sorted subsequences**
- Processor $p-1$ determines the splitters
 - ♦ It samples each **sorted subsequence**, taking every $(kn/(p^2s))^{\text{th}}$ element ($1 \leq k \leq s-1$), where $p \leq s \leq n/p^2$
 - ♦ Merges the sampled sequences, and collects $p-1$ regularly spaced splitters
 - ♦ Broadcasts the splitters to all processors
- Processors route (exchange) **sorted subsequences** according to the splitters (transpose)
- The data are unshuffled

Performance

- Assuming $n \geq p^3$ and $p \leq s \leq n/p^2$
- Running time is $\approx O((n/p) \lg n)$
- With high probability ...
no processor holds more than $(n/p + n/s - p)$ elements
- Duplicates d do not impact performance unless $d = O(n/p)$
- Tradeoff: increasing s ...
 - ♦ Spreads the final distribution more evenly over the processors
 - ♦ Increases the cost of determining the splitters
- For some inputs, communication patterns can be highly irregular with some pairs of processors communicating more heavily than others
- This imbalance degrades communication performance

The collective calls

- Processes transmit varying amounts of information to the other processes
- This is an `MPI_Alltoallv`
(`SKeys`, `send_counts`, `send_displace`, `MPI_INT`,
`RKeys`, `recv_counts`, `recv_displace`, `MPI_INT`,
`MPI_COMM_WORLD`)
- Prior to making this call, all processes must cooperate to determine how much information they will exchange
 - ♦ The *send list* describes the number of keys to send to each process k, and the offset in the local array
 - ♦ The *receive list* describes the number of incoming keys for each process k and the offset into the local array

Determining the send and receive lists

- After sorting, each process scans its local keys from left to right, marking where the splitters divide the keys, in terms of **send counts**
- Perform an all to all to transpose these send counts into **receive counts**

```
MPI_Alltoall(send_counts, 1, MPI_INT,  
             recv_counts, 1, MPI_INT, MPI_COMM_WORLD)
```

- A simple loop determines the displacements

```
for (p=1; p < nodes; p++){  
    s_displ[p] = s_displ[p-1] + send_counts[p-1];  
    r_displ[p] = r_displ[p-1] + recv_counts[p-1];  
}
```

Fin